

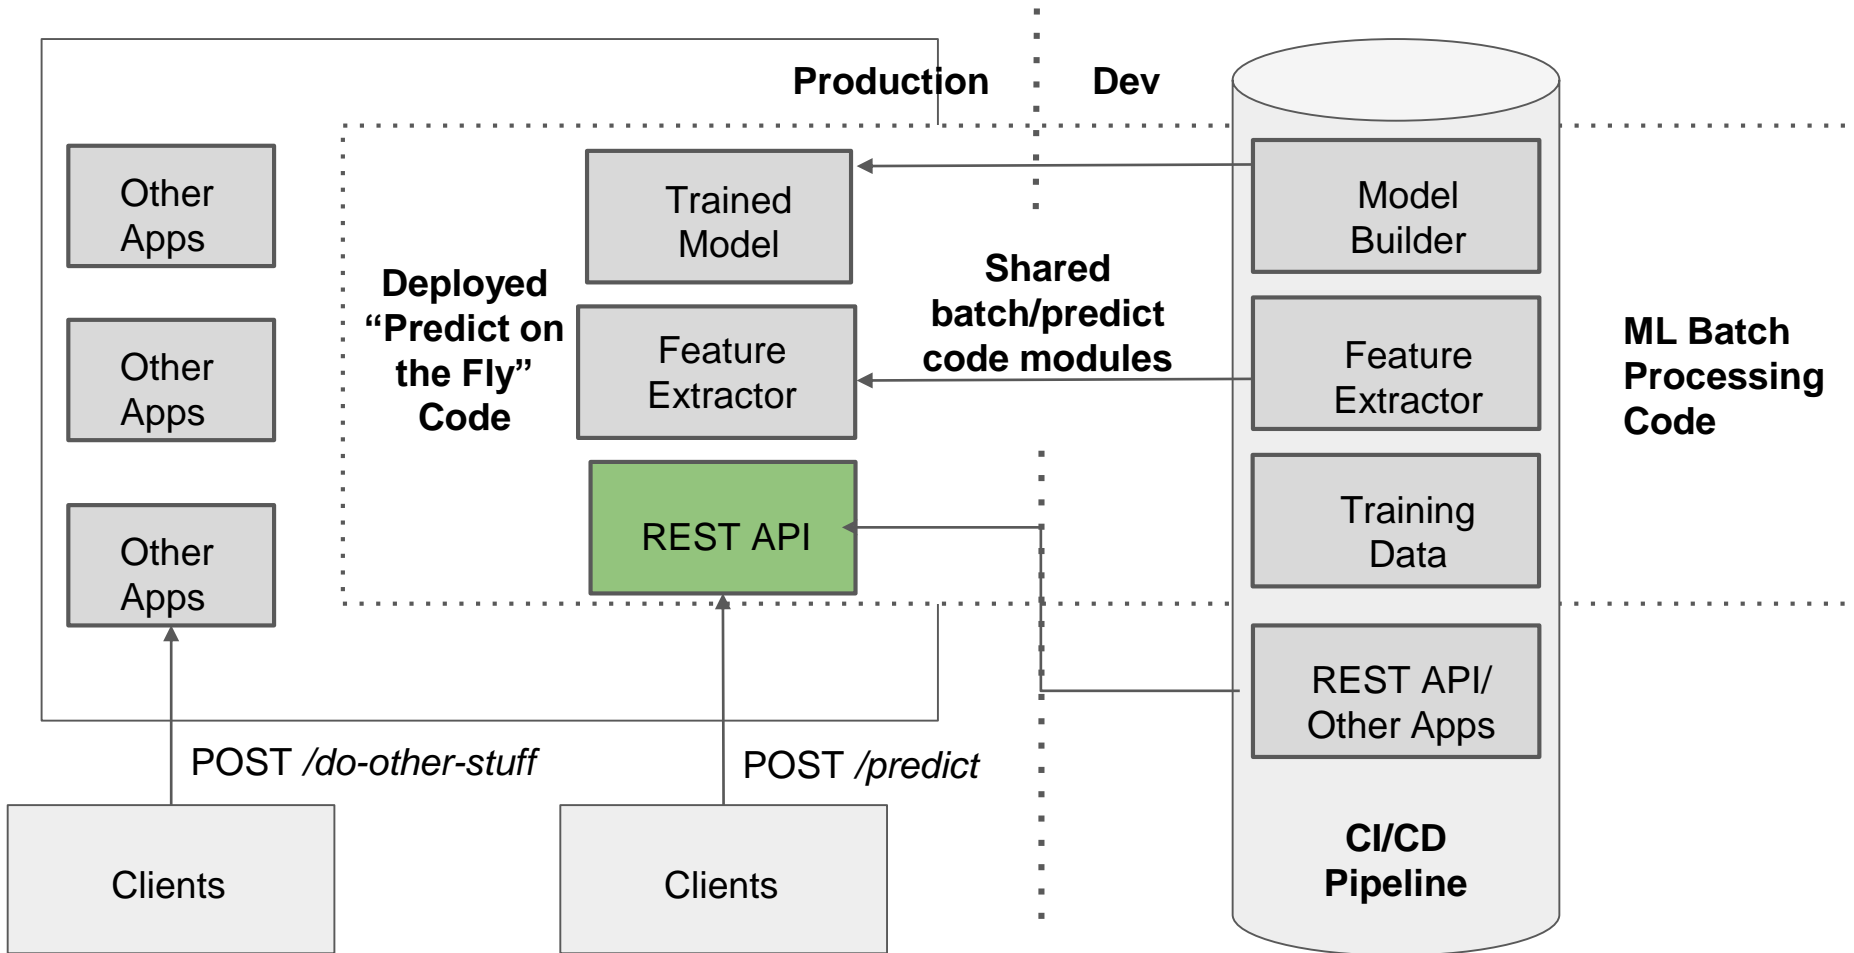
Section 7: Serving the Model

Introduction



What we will be Covering in this Section

Section 7



What is a REST API?

A Representational State Transfer (REST)
Application Programming Interface (API):

- Surprisingly loaded question!
- Very simply, a server will transfer to the client a representation of the state of a requested resource.
- Popular due to its simplicity
- See the lecture notes for details and theory



Why Does This Matter?

Serving our model via REST API allows us to:

- Serve predictions on the fly to multiple clients
- Decouple our model development from the client facing layer
- Potentially combine multiple models at different API endpoints
- Scale by adding more instances of the application behind a load balancer



Flask Overview



We will build our API using the flask micro framework:

- Popular choice for Python microservices (40k stars on github)
- Lightweight and flexible
- Alternatives include: Django, Pyramid, Bottle, Sanic, Tornado, API Star amongst many others.

Let's Get Started!

See you in the next section